

# Software Development Intern Assignment

## About Us:

ReachInbox is transforming cold outreach with our revolutionary AI-driven platform. Our all-in-one solution empowers businesses to effortlessly find, enrich, and engage high-intent leads through multi-channel outreach on Twitter, LinkedIn, email, and phone. With just a single prompt, ReachInbox springs into action, prospecting and verifying leads, crafting personalized sequences, and notifying businesses of responsive prospects. Imagine being part of an AI-powered growth team that consistently generates top-tier leads. ReachInbox is more than a tool; it's your growth partner.

We are looking for passionate and innovative individuals to join our team and help us continue to redefine the future of lead generation and business growth.

## ReachInbox Hiring Assignment – Full-stack Email Job Scheduler

### Problem Statement

At ReachInbox, a huge part of our system is **reliable scheduling and sending of emails at scale**.

Your task is to build a **production-grade email scheduler service + dashboard** that:

- Accepts **email send requests** via APIs
- Schedules them to be sent at a **specific time**
- Uses **BullMQ + Redis** as a persistent job scheduler (**no cron jobs**)
- Sends emails using fake SMTP via **Ethereal Email**
- Survives **server restarts** without restarting from scratch or losing jobs
- Exposes a **frontend dashboard** to:
  - Schedule new emails
  - View scheduled emails
  - View sent emails

Think of it as a tiny slice of what ReachInbox does under the hood.

---



# Tech Requirements

You **must** use:

## Backend

- Language: **TypeScript**
- Framework: **Express.js**
- Queue: **BullMQ** (backed by **Redis**)
- Database: **MySQL or PostgreSQL** (ORM or query builder is your choice)
- SMTP: **Ethereal Email** (fake SMTP for testing)

## Frontend

- **React.js or Next.js**
- **Tailwind CSS** (or any modern CSS library)
- **TypeScript** strongly preferred

## Infra

- Redis and DB can be run via **Docker** (recommended but not mandatory)

We care about how you **structure a real backend, wire queues correctly, and build a clean frontend.**

---



## Backend Requirements

### 1 Core Scheduler Behavior

Your backend must:

- Accept email scheduling requests via API
- Store them in a **relational DB** (MySQL/Postgres)
- Schedule them using **BullMQ delayed jobs** (or a custom Redis/DB-based scheduler) – **no cron**
- Send emails from multiple senders via **Ethereal Email (SMTP)**
- Persist state so that:
  - If the server restarts, **future emails are still sent at the correct time**
  - Emails are **not duplicated** or restarted from scratch

### 2 Throughput, Rate Limiting & Concurrency (Required)

Your scheduler should behave like a **real-world email system under load**.

You **must** support:

## Worker Concurrency

- Configure your BullMQ worker(s) with a **configurable concurrency** level .
- Implementation must be safe when multiple jobs run in parallel:

## Delay Between Each Email

- There must be a **minimum delay between individual email sends** (to mimic provider throttling).
- You can:
  - Use **BullMQ's limiter** options, or
  - Add a custom delay in the worker logic.
- Document in the README what delay you chose (e.g. “min 2 seconds between sends”).

## Emails Per Hour (Rate Limiting)

- Implement a **rate limit** on the number of emails sent per hour:
  - Either **global** (e.g. `MAX_EMAILS_PER_HOUR=200`), or
  - **Per-sender / per-tenant** (e.g. `MAX_EMAILS_PER_HOUR_PER_SENDER`), you will have to support multiple senders.
- The **limit values must be configurable** via env/config (no hardcoding).
- Rate-limiting logic must be **safe across multiple workers / instances**:
  - Use **Redis or DB** backed counters (e.g. keyed by `hour_window + sender`),
  - **Do not** rely only on in-memory counts.
- When the hourly limit is reached:
  - **Do not drop** or permanently fail jobs.
  - Jobs should be **delayed / rescheduled** into the **next available hour window** while preserving order as much as possible.
- Explain in the README how you enforce this:
  - e.g. BullMQ limiter, Redis counters, custom logic, trade-offs, etc.

## Behavior Under Load

Your design should clearly define behavior when:

- **1000+ emails** are scheduled for roughly the same time.
- The **rate limit** would be exceeded.

You don't need to actually send thousands via Ethereal, but your logic should handle it.

## 3 Hard Constraints (Important)

These are **non-negotiable**:

- **✗ Do NOT use cron jobs**
    - No OS-level cron (`crontab`, etc.)
    - No Node cron libraries (`node-cron`, `agenda`, etc.)
  - **✓ Scheduling must be done using:**
    - **BullMQ delayed jobs**, or
    - A **custom scheduler** that uses Redis/DB to track next run times — but still **not cron**.
  - **✓ The system must be persistent:**
    - After a restart:
      - Future scheduled emails **still send at the right time**
      - Emails are **not re-sent** or restarted from Day 1
  - **✗ Same email queues should not be sent more than once. Maintain idempotency.**
- 



## Frontend Requirements

You will build a **frontend that matches the provided [Figma](#)** as closely as possible and talks to your backend APIs.

### [Figma Link](#)

## 1 Google Login (Required)

- Implement **real Google OAuth login** (no mock).
- After login, redirect the user to the **dashboard**.
- Show in the top header:
  - User's **name**
  - **Email**
  - **Avatar**
- Provide a simple **Logout** option.

## 2 Main Dashboard

After login, show the main UI with:

- **Top header** (user info + logout).
- Tabs/sections:
  - **Scheduled Emails**
  - **Sent Emails**

- A primary “**Compose New Email**” button.

Layout and styling should closely follow the **Figma design**.

### 3 Compose New Email

User must be able to:

- Enter:
  - **Subject**
  - **Body**
- **Upload** a CSV/text file of email leads.
  - Parse and show the **number of email addresses** detected.
- Set:
  - **Start time** (when scheduling begins)
  - **Delay between emails**
  - **Hourly limit**
- Click **Schedule** to send data to the **backend schedule API**.

This can be a **modal** or **separate page**, depending on the Figma.

### 4 Scheduled Emails

Show a clean table/list with:

- **Email**
- **Subject**
- **Scheduled time**
- **Status**

Include:

- **Loading states**
- **Empty state** when there are no scheduled emails

### 5 Sent Emails

Show a table/list with:

- **Email**
- **Subject**
- **Sent time**
- **Status** (**sent** / **failed**)

Include:

- **Loading states**
- **Empty state** when there are no sent emails

## 6 Frontend Code Quality

We expect:

- Clean **folder structure**
  - Reusable **UI components** (buttons, inputs, tables, modals, etc.)
  - **DRY** code (avoid duplication)
  - Proper **TypeScript** usage:
    - Types/interfaces for API responses & props
  - Good UX:
    - Loading indicators
    - Empty states
    - Error handling (basic messages/toasts)
- 



## Submission Guidelines

1. **Create a private GitHub repository** (monorepo or separate folders for backend & frontend is fine).
2. **Grant access** to user: [Mitrajit](#)
3. Add a **README** that includes:
  - How to run **backend** (Express, Redis, DB, BullMQ worker)
  - How to run **frontend**
  - How to set up **Ethereal Email** and env variables
  - Architecture overview:
    - How scheduling works
    - How persistence on restart is handled
    - How rate limiting & concurrency are implemented
  - List of **features implemented**, mapped to:
    - Backend: scheduler, persistence, rate limiting, concurrency
    - Frontend: login, dashboard, compose, tables, etc.
4. Add a **short demo video (max 5 minutes)**:
  - Show creating scheduled emails (from frontend or Postman)
  - Show the dashboard with **Scheduled** and **Sent** emails
  - Show a **restart scenario**:
    - Stop server → start again → future emails still send
  - (Bonus) Briefly demonstrate how **rate limiting / delay** behaves under load
5. Note any **assumptions, shortcuts, or trade-offs** you made.

<aside>💡 \*\*Fill this form with relevant links and details - Assignment Submission - <https://forms.gle/SRv2FoT913n2B2yg6>

</aside>

### **Deadline for Task Completion:**

You have a maximum of 48 hours to complete the task. Receiving this assignment means you're already ahead of many candidates. Good luck!

**Note:** Do not submit a plagiarized assignment. All GitHub code will be thoroughly reviewed, and any evidence of plagiarism will result in the assignment being rejected.

Thank you!